DELTA: Edge-Centric Dual Attention for Compositional Reasoning on Knowledge Graphs

Benjamin Brown
Independent Researcher
bdbrown4@eagles.usi.edu

Abstract

Graph neural networks treat edges as passive conduits for message passing between nodes, limiting their ability to reason about *relationships between relationships*—the compositional structure underlying multi-hop inference in knowledge graphs. We introduce **DELTA** (**D**ual **E**dge-**L**inked **T**ransformer **A**rchitecture), which elevates edges to first-class computational entities via *edge-to-edge attention*: edges attend to structurally adjacent edges through a multi-hop edge adjacency matrix, enabling direct relational composition without node-mediated bottlenecks. On three synthetic structural benchmarks (edge-type relation classification, noise robustness at 80% corruption, and compositional relation discovery), evaluated across 5 seeds on an H100, DELTA outperforms GraphGPS and GRIT by $+0.573$, $+0.290$, and $+0.095$ accuracy respectively—providing the primary cross-task evidence that edge-to-edge attention captures a structurally distinct reasoning capacity. On the complete 14,541-entity FB15k-237 graph (mean degree ≈4.1), a 3-seed hop-depth ablation confirms that 2-hop edge adjacency improves 3-hop chain query MRR by $+0.012$ over 1-hop adjacency (0.218 vs. 0.206), validating the core mechanism at realistic scale. As an exploratory characterization on a top-500-entity degree-based subgraph (mean degree ≈19.7), DELTA-Matched (158K parameters) is the only model among 7 tested with a depth-monotonic MRR improvement from 2-hop (0.758) to 5-hop (0.790); GraphGPS degrades from 0.754 to 0.690, and the advantage reaches $+0.100$ MRR at 5-hop. Cross-seed validation confirms the subgraph trend (0.742 ± 0.009 vs. 0.713 ± 0.007 at 3-hop; non-overlapping CIs). However, a direct mechanism ablation on this dense subgraph finds 1-hop, 2-hop, and node-only conditions statistically indistinguishable, consistent with high link density saturating 1-hop adjacency coverage. An emergent edge–node temperature asymmetry arises during training: edge attention heads converge to sharper distributions while node heads prefer softer averaging, a pattern discovered autonomously via gradient descent. DELTA achieves this with 31% fewer parameters than GraphGPS and faster per-query inference (0.8–0.9×), at the cost of a one-time encoding overhead from edge adjacency construction. The central open question for future work is whether a SOTA-performing DELTA variant—one incorporating specialized KGC training strategies at full-graph scale—preserves the depth-monotonic compositional advantage observed here.

1 Introduction

Knowledge graph completion (KGC) requires inferring missing links from observed relational structure. While standard link prediction evaluates single-hop tail prediction ("1p"), real-world reasoning often demands *compositional* multi-hop inference: chaining relations such as `born_in` ∘ `located_in` to derive `nationality`. This compositional capacity is precisely where existing architectures falter.

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

Embedding methods (TransE [1], DistMult [2], RotatE [3]) represent entities and relations as vectors without structural inductive bias. They excel at memorizing local link statistics but collapse on multi-hop chains: DistMult's MRR drops from 0.728 (2-hop) to 0.457 (5-hop) on our top-500-entity subgraph benchmark. Message-passing GNNs (CompGCN [7], R-GCN [6]) and graph transformers (GraphGPS [10]) propagate information through nodes, treating edges as scalar weights or simple feature carriers. This design creates an information bottleneck: to reason about relationships between relationships, information must route through intermediate nodes, accumulating compression loss at each hop.

We propose **DELTA**, an architecture that makes edges first-class computational entities. DELTA's key mechanism is *edge-to-edge attention via multi-hop edge adjacency*: two edges that share a node are "adjacent" in edge space, and at 2 hops, an edge can attend to edges two steps away. This creates a direct channel for relational composition—`born_in` can directly attend to `located_in` without routing through the intermediate city node—enabling transitive inference as a natural consequence of the attention structure rather than an emergent property that must be learned.

DELTA runs node attention and edge attention as parallel streams, reconciled through a bidirectional bridge that ensures cross-stream coherence. This *dual parallel attention* converges 2.7$\times$ faster than sequential alternatives while maintaining the same representational capacity.

Contributions:

1. **Edge-to-edge attention architecture.** We introduce dual parallel attention with multi-hop edge adjacency, the first GNN mechanism that enables direct edge-to-edge reasoning without node-mediated bottlenecks (§3).
2. **Cross-task synthetic validation.** On three synthetic structural benchmarks—edge-type relation classification, edge classification under 80% noise corruption, and compositional relation discovery (4 base + 3 derived rules)—DELTA achieves 0.880, 1.000, and 1.000 mean accuracy (5 seeds, H100), vs. GraphGPS 0.293/0.697/0.905 and GRIT 0.307/0.710/0.893. These tasks are entirely distinct from knowledge graph completion and provide cross-task evidence that edge-to-edge attention captures a general structural reasoning capacity (§4.2).
3. **Full-scale validation on complete FB15k-237.** A direct 3-seed hop-depth ablation on the full 14,541-entity graph confirms that 2-hop edge adjacency improves 3-hop chain query MRR by $+0.012$ over 1-hop (0.218 vs. 0.206) on the sparse graph (mean degree $\approx$4.1) where the mechanism is most relevant (§4.5).
4. **Exploratory depth characterization on degree-based subgraph.** On a 500-entity degree-based subgraph of FB15k-237, DELTA-Matched exhibits a depth-monotonic MRR improvement from 2-hop to 5-hop that no other tested model shares, with the advantage over GraphGPS accelerating from $+0.004$ to $+0.100$ MRR. A direct mechanism ablation on this subgraph finds all hop-depth conditions statistically indistinguishable (consistent with high density saturating 1-hop adjacency), so these results characterize emergent behavioral depth-dependence rather than isolating the mechanism (§4.3).
5. **Emergent attention temperature asymmetry.** We observe that edge and node attention heads develop fundamentally different sharpness preferences during training: edge heads drift toward higher temperature (sharper) while node heads drift lower (softer). This asymmetry is consistent across 20+ configurations and 9 experimental phases, suggesting distinct computational roles for the two attention streams (§4.7).
6. **Scalability characterization.** We evaluate DELTA at scales from N=500 to N=14,541 entities, identify attention dilution as the primary scaling bottleneck, and validate top-k sparse attention (k=128) as a lossless approximation that preserves full-softmax quality (§4.8).

2 Related Work

Knowledge graph embeddings. TransE [1] models relations as translations in embedding space; DistMult [2] uses bilinear scoring; RotatE [3] models relations as rotations in complex space; ComplEx [4] extends to complex-valued embeddings. These methods achieve strong single-hop link prediction but lack structural inductive bias for multi-hop reasoning. PairRE [5] and RotatE achieve

MRR $\sim$0.34 on full FB15k-237, but rely entirely on embedding geometry rather than graph structure for relational composition.

Message-passing GNNs for KGs. R-GCN [6] introduces relation-specific weight matrices for typed edges. CompGCN [7] jointly embeds nodes and relations through composition operations (subtraction, multiplication, circular correlation). Both treat edges as typed message channels rather than computational entities—edge representations are derived from relation type embeddings, not learned through edge-to-edge interaction. NBFNet [8] and RED-GNN [9] address multi-hop reasoning via path-based message passing but still route all information through nodes. A*Net [21] scales path-based reasoning to large KGs via priority-based entity selection. These path-based methods are a key baseline comparison target for future work; our primary comparison here is against general graph transformer architectures (GraphGPS, GRIT).

Graph transformers. GraphGPS [10] combines local MPNN message passing with global transformer attention, using positional/structural encodings to inject graph awareness into the attention mechanism. GRIT [11] uses random-walk structural encodings with graph transformers. Both architectures provide expressive node-level representations but lack mechanisms for direct edge-to-edge interaction. Exphormer [12] addresses scalability via sparse attention patterns but does not introduce edge-level attention.

Edge-aware architectures. Line graphs [13] and their GNN variants operate on the dual graph where edges become nodes, enabling edge-level message passing. However, line graph construction is $O(E \cdot d_{\max})$ and the resulting graph can be much larger than the original. DELTA's edge adjacency matrix provides similar connectivity at $O(E^{0.97})$ via sparse COO operations, with the crucial difference that DELTA maintains both node and edge representations simultaneously through dual attention. EHGNN [14] incorporates edge features into heterogeneous GNNs but does not implement edge-to-edge attention.

Multi-hop reasoning on KGs. BetaE [15] and ConE [16] embed logical queries as geometric regions for answering complex queries. GQE [17] and Q2B [18] project queries into embedding spaces. CQD [22] answers complex queries by continuous optimization over link predictor scores. GNN-QE [23] combines GNN encoders with fuzzy logic for complex query answering. These methods are designed specifically for complex query answering rather than general KGC. Our multi-hop evaluation uses chain queries (2p–5p) with soft entity traversal—a complementary protocol that tests compositional capacity within standard KGC training, and is not directly comparable to BetaE/GQE query benchmark numbers.

3 Architecture

DELTA processes a knowledge graph $G = (V, E, R)$ with entity set V, edge set $E \subseteq V \times R \times V$, and relation set R. Each entity v_i has a learned embedding $\mathbf{h}_i \in \mathbb{R}^{d_n}$ and each edge (v_i, r_k, v_j) has a learned embedding $\mathbf{e}_{ij} \in \mathbb{R}^{d_e}$.

3.1 Multi-Hop Edge Adjacency

The foundation of DELTA is the *edge adjacency matrix* $\mathbf{A}_E$, which encodes structural connectivity between edges. Two edges $e_a = (v_i, r, v_j)$ and $e_b = (v_k, r', v_l)$ are 1-hop adjacent if they share an endpoint: $\{v_i, v_j\} \cap \{v_k, v_l\} \neq \emptyset$. The 2-hop edge adjacency extends this via the incidence structure:

$$\mathbf{A}_E^{(2)} = (\mathbf{B}^\top \mathbf{B})^2 - \text{diag} \tag{1}$$

where $\mathbf{B} \in \{0, 1\}^{|V| \times |E|}$ is the node-edge incidence matrix. This is computed efficiently using sparse COO tensor operations via `torch_sparse.spspmm`, achieving $O(E^{0.97})$ scaling (validated empirically from 50 to 2,500 edges). At 2 hops, an edge can "see" edges two steps away, enabling direct reasoning about relational chains such as `born_in` $\rightarrow$ `located_in` without intermediate node representations.

We validate the role of 2-hop edge adjacency in §4.5: on a synthetic transitive relation benchmark, 2-hop achieves 100% accuracy on derived relations vs. 61.1% for 1-hop. On FB15k-237 multi-hop queries, a direct ablation finds all three adjacency configurations perform comparably (Table 7); we discuss this result and its implications in §4.5.

3.2 Edge-to-Edge Attention

Given the edge adjacency $\mathbf{A}_E^{(2)}$ stored as a sparse index $[2, E_{\text{adj}}]$, we compute multi-head edge attention. For edge e_i attending to adjacent edge e_j:

$$\alpha_{ij}^{(h)} = \frac{\tau^{(h)} \cdot (\mathbf{W}_Q^{(h)}\mathbf{e}_i)^\top (\mathbf{W}_K^{(h)}\mathbf{e}_j)}{\sqrt{d_h}} + \mathbf{W}_{\text{ctx}}[\mathbf{h}_{s(j)} \| \mathbf{h}_{t(j)}]_h \tag{2}$$

where $\tau^{(h)} = \exp(\theta^{(h)})$ is a learnable per-head temperature (always positive). Note that $\tau^{(h)}$ *multiplies* the dot product (not divides): larger τ produces sharper, more selective attention—the inverse of the conventional divide-by-temperature softmax formulation. $s(j)$ and $t(j)$ are the source and target nodes of edge e_j, $\|$ denotes concatenation, and $d_h = d_e/H$ is the head dimension. The node context bias $\mathbf{W}_{\text{ctx}} \in \mathbb{R}^{2d_n \times H}$ grounds edge attention in the endpoint identities of the attending edges.

Attention weights are computed via scatter-softmax over the sparse adjacency structure (normalizing per target edge), followed by weighted aggregation:

$$\mathbf{e}_i' = \text{LayerNorm}\left(\mathbf{e}_i + \mathbf{W}_O \sum_{j \in \mathcal{N}_E(i)} \text{softmax}(\alpha_{ij}) \cdot \mathbf{W}_V \mathbf{e}_j\right) \tag{3}$$

Top-k sparse attention. For large graphs where $|\mathcal{N}_E(i)|$ becomes prohibitive, we retain only the k highest-scoring neighbors per target edge (based on mean attention across heads), then renormalize. At N=5,000 entities with 63M adjacency pairs, k=128 preserves full-softmax quality exactly (test MRR 0.2472 vs. 0.2457 for full softmax), while k=64 degrades by -5.6%.

3.3 Node Attention

Node attention follows a GAT-style [20] design where each node attends to its neighbors, with edge features providing a per-head bias:

$$\beta_{ij}^{(h)} = \frac{\tau_n^{(h)} \cdot (\mathbf{W}_Q^{(h)}\mathbf{h}_i)^\top (\mathbf{W}_K^{(h)}\mathbf{h}_j)}{\sqrt{d_h}} + \mathbf{W}_{\text{edge}}(\mathbf{e}_{ij})_h \tag{4}$$

Node attention uses independent learnable temperatures $\tau_n^{(h)}$ per head, enabling the model to discover optimal sharpness for each attention stream independently.

3.4 Dual Parallel Attention

Both attention mechanisms operate on the *same* input state simultaneously:

$$\mathbf{H}' = \text{NodeAttn}(\mathbf{H}, \mathbf{E}, G), \qquad \mathbf{E}' = \text{EdgeAttn}(\mathbf{E}, \mathbf{H}, \mathbf{A}_E) \tag{5}$$

This parallel design ensures that neither stream's output biases the other's computation within a layer, and converges 2.7× faster than sequential alternatives.

3.5 Reconciliation Bridge

After parallel attention, a bidirectional bridge ensures cross-stream coherence:

$$\hat{\mathbf{e}}_{ij} = \text{LN}(\mathbf{e}'_{ij} + \mathbf{W}_1[\mathbf{e}'_{ij} \| \mathbf{h}'_{s(ij)} \| \mathbf{h}'_{t(ij)}]) \quad \text{(edges absorb node context)} \tag{6}$$

$$\hat{\mathbf{h}}_i = \text{LN}(\mathbf{h}'_i + \mathbf{W}_2[\mathbf{h}'_i \| \bar{\mathbf{e}}_i]) \quad \text{(nodes absorb edge context)} \tag{7}$$

where $\bar{\mathbf{e}}_i = \text{mean}(\{\hat{\mathbf{e}}_{ij} : j \in \mathcal{N}(i)\})$ aggregates incident updated edge features. Both updates use residual connections with LayerNorm. Edges update first (absorbing node context), then nodes update (absorbing the already node-informed edge features), creating a sequential cascade within the bridge.

3.6 Link Prediction Head

For link prediction, we use a DistMult [2] decoder:

$$s(v_i, r_k, v_j) = \mathbf{h}_i^\top \text{diag}(\mathbf{r}_k)\mathbf{h}_j \tag{8}$$

where $\mathbf{r}_k$ is a learned relation embedding. Training uses 1-vs-all binary cross-entropy with label smoothing (0.1), following standard KGC practice.

3.7 Multi-Hop Scoring

For n-hop path queries $(v_0, r_1, ?, r_2, ?, \ldots, r_n, ?)$, we use *soft entity traversal*: at each intermediate hop, score all entities via the DistMult decoder, apply softmax with temperature $\tau = 1.0$ to obtain a probability distribution over entities, take the weighted average entity embedding, and pass it to the next hop. This provides a differentiable approximation to hard entity traversal while remaining faithful to the compositional structure of the query.

4 Experiments

4.1 Setup

Dataset. We use FB15k-237 [19], the standard filtered subset of Freebase with 14,541 entities, 237 relations, and 310,116 triples. Our compositional reasoning and mechanism analysis experiments (§4.3–4.7) use a **top-500 degree subgraph** (494 entities, 207 relations, ∼9,700 train triples, mean degree ≈19.7) with full filtered evaluation. This degree-biased subgraph enables rapid multi-configuration ablations across many seeds and depth levels. Scaling experiments (§4.8) extend to N=2,000 and N=5,000 entity subgraphs. The hop-depth ablation (§4.5) additionally reports a full-scale evaluation on the complete 14,541-entity graph (mean degree ≈4.1, 310,116 triples) to validate key findings at standard benchmark scale.

Models. We compare seven configurations:

- **DELTA-Matched** (158K params): 1 DELTA layer, d_n=64, d_e=32, 4 heads. Our primary model, designed for parameter-fair comparison.
- **DELTA-Full** (293K params): 3 DELTA layers.
- **SelfBootstrapHybrid** (381K params): A FixedChain DELTA layer bootstraps graph construction before 3 DELTA processing layers.
- **GraphGPS** [10] (228K params): Local MPNN + global transformer attention.
- **GRIT** [11] (197K params): Random-walk structural encodings + graph transformer.
- **DistMult** [2] (47K params): Pure embedding baseline (no GNN).
- **SelfBootstrap** (299K params): DELTA bootstraps its own graph construction.

Training. All models use Adam optimizer (lr $= 0.003$), batch size 4,096, 1-vs-all BCE loss with label smoothing 0.1, and early stopping with patience 10. Evaluation uses standard filtered MRR and Hits@k metrics. Multi-hop queries use 10% DropEdge for DELTA (validated in §4.3).

Table 1: Phase 11 — transitive relation classification accuracy by adjacency configuration (100-entity synthetic graph, single seed). Only 2-hop edge adjacency achieves perfect accuracy on derived (composed) relation instances.

Config	Transitive Acc.	Δ vs. 2-hop
2-hop (DELTA)	**1.000**	—
Node-GNN (no edge attn)	0.833	-0.167
1-hop	0.611	-0.389
MLP (no structure)	0.375	-0.625

Multi-hop query generation. Chain queries are generated by composing training edges to reach test edges: a 2-hop query chains one train edge to one test edge, a 3-hop query chains two train edges to one test edge, etc. All queries undergo strict leakage prevention: no 1-hop shortcuts between anchor and answer in the training graph, no duplicate queries, and automated audit verification before evaluation. Query counts: 1p=486, 2p=5,382, 3p=10,000, 4p=10,000, 5p=10,000.

4.2 Synthetic Structural Benchmarks

Before evaluating on the FB15k-237 knowledge graph, we validate the edge-to-edge attention mechanism on synthetic structural tasks that isolate compositional reasoning without confounds from embedding quality or training scale.

Transitive edge adjacency (Phase 11). On a 100-entity synthetic knowledge graph with explicitly transitive relations, we ablate edge adjacency depth (Table 1). 2-hop edge adjacency achieves **100%** accuracy on derived (transitive) relation instances, vs. 1-hop adjacency at 61.1%, a node GNN at 83.3%, and an MLP at 37.5%. 2-hop adjacency is necessary and sufficient for transitive relational composition. This result, alongside the full-graph ablation ($+0.012$ on 3p MRR), constitutes the primary mechanistic evidence for the hop-depth claim, since the dense-subgraph mechanism ablation (Phase 66, Table 7) finds all conditions statistically indistinguishable.

Three-task benchmark (Phase 34). We compare DELTA-Matched, GraphGPS, and GRIT on three synthetic structural tasks (5 seeds $\times$ 500 epochs, H100; Table 2):

- **Edge-type classification:** A synthetic KG with 6 relation patterns (8 instances each); classify each edge by its relation type. DELTA achieves 0.880 vs. GraphGPS 0.293 and GRIT 0.307 (advantage $+0.573$). Edge-to-edge attention learns discriminative edge-type representations that node-centric message passing cannot extract—standard graph transformers have no native edge-classification head, and their node representations lose the information needed to distinguish edge types.
- **Noise robustness:** Repeat edge-type classification with 80% of edges randomly corrupted. DELTA maintains 1.000 accuracy while GraphGPS degrades to 0.697 and GRIT to 0.710. DELTA's edge-level representation is structurally grounded and robust to corrupted neighborhood signals.
- **Compositional relation discovery:** A KG with 4 base relation types and 3 derived (compositional) rules; models must classify edges including derived instances produced by relational composition. DELTA achieves 1.000 vs. GraphGPS 0.905 ($+0.095$), directly demonstrating that 2-hop edge adjacency enables compositional rule discovery without supervision.

The cross-task pattern is consistent: DELTA's largest advantage is on the edge-type classification task ($+0.573$), where node-centric architectures are structurally mismatched. The advantage is smallest on path composition ($+0.095$), where multi-hop node-level message passing provides a partial substitute. Together with the transitive adjacency ablation, these results show that DELTA's mechanism is not FB15k-237-specific—edge-to-edge attention provides a general structural advantage wherever edge identity matters.

Probe architecture note. Standard graph transformers (GraphGPS, GRIT) have no native edge-classification head; the comparison equips all models with the same linear probe, but probe inputs differ by necessity: DELTA's probe operates on learned edge representations, while GraphGPS/GRIT

Table 2: Synthetic structural benchmarks (Phase 34): mean accuracy across 5 seeds (500 epochs, H100), compared at matched parameter budgets. Advantage row shows DELTA vs. the stronger of GraphGPS/GRIT.

Model	Params	Edge Cls.	Noise @0.8	Path Comp.
DELTA-Matched	158K	**0.880**	**1.000**	**1.000**
GraphGPS	228K	0.293	0.697	0.905
GRIT	197K	0.307	0.710	0.893
DELTA advantage	—	**+0.573**	**+0.290**	**+0.095**

Table 3: Multi-hop MRR by reasoning depth (single seed, 0% DropEdge). DELTA-Matched is the only model with higher MRR at 5-hop than 2-hop (net $+0.032$). All baselines show net degradation from 2-hop to 5-hop.

Model	Params	1p	2p	3p	4p	5p
DELTA-Matched	158K	0.541	**0.758**	**0.753**	**0.767**	**0.790**
GraphGPS	228K	**0.523**	0.754	0.727	0.701	0.690
DistMult	47K	0.494	0.728	0.583	0.511	0.457
DELTA – GraphGPS gap		+0.018	+0.004	+0.026	+0.066	**+0.100**

probes use mean-pooled endpoint node embeddings. This architectural asymmetry may inflate DELTA's advantage on the edge-type classification task specifically. The path composition result ($+0.095$), where all models predict via the same endpoint protocol, provides a cleaner cross-architecture comparison. Future work should equip GraphGPS/GRIT with matched learned edge projections and measure the residual gap.

4.3 Compositional Reasoning

We characterize depth-dependent performance on the FB15k-237 degree-based subgraph. DELTA-Matched is the only model among 7 tested that shows a depth-monotonic MRR improvement from 2-hop to 5-hop; GraphGPS and all other baselines degrade. Readers should note that a direct mechanism ablation on the same subgraph (Table 7) finds 1-hop, 2-hop, and node-only conditions statistically indistinguishable—the depth-monotonic pattern is a behavioral property of the model on this graph, not an isolated consequence of the hop-depth setting. The full-graph ablation (§4.5) provides the mechanistic evidence that 2-hop adjacency adds discriminative value on sparse graphs.

The depth advantage accelerates. Table 3 presents the central result. DELTA-Matched's lead over GraphGPS grows from $+0.004$ (2p) to $+0.100$ (5p)—a $25\times$ amplification. Critically, DELTA's 5p MRR (0.790) *exceeds* its own 2p MRR (0.758). No other model exhibits this property: GraphGPS loses -0.064 MRR from 2p→5p, while DistMult loses -0.271.

Consistent improvement under regularization. Table 4 shows the full 7-model comparison at seed=1 with 0% DropEdge. DELTA-Matched shows a small net gain from 2p (0.733) to 3p (0.738), while every other model degrades; the same configuration shows a brief 2p→3p dip in Table 3 before recovering to 5p. The multi-seed evaluation (Table 5) with 10% DropEdge confirms the advantage is robust: DELTA @10% DropEdge improves from 2p (0.730 ± 0.011) to 3p (0.742 ± 0.009), while GraphGPS degrades ($0.727 \rightarrow 0.713$). The 2p→3p improvement is consistent across seeds under regularization. This suggests the advantage stems from parameter efficiency combined with edge-to-edge attention: larger models overfit to 1-hop statistics, while DELTA-Matched is forced to learn genuinely compositional representations.

Multi-seed validation. To confirm statistical robustness, we run DELTA-Matched (with 10% DropEdge) and GraphGPS across 3 seeds (Table 5). At 3-hop, DELTA achieves 0.742 ± 0.009 vs. GraphGPS 0.713 ± 0.007. The confidence intervals do not overlap: DELTA's worst seed (0.731) exceeds GraphGPS's best (0.722).

Table 4: Complete multi-hop results (7 models, seed=1). Only DELTA-Matched improves from 2p→3p.

Model	Params	1p	2p	3p	3p H@10	2p→3p
DELTA-Matched	158K	0.533	**0.733**	**0.738**	**0.873**	**+0.005**
SBHybrid	381K	**0.541**	0.722	0.695	0.833	-0.027
GraphGPS	228K	0.549	0.718	0.697	0.850	-0.021
DELTA-Full	293K	0.524	0.711	0.692	0.836	-0.020
SelfBootstrap	299K	0.512	0.712	0.686	0.829	-0.026
GRIT	197K	0.460	0.712	0.644	0.783	-0.068
DistMult	47K	0.532	0.715	0.566	0.710	-0.150

Table 5: Multi-seed 3-hop MRR ($\pm$ std, 3 seeds). Non-overlapping confidence intervals.

Config	Params	1p	2p	3p	2p→3p
DELTA @10% DE	158K	0.543 ± 0.006	0.730 ± 0.011	$\mathbf{0.742 \pm 0.009}$	+0.012
GraphGPS	228K	0.529 ± 0.009	0.727 ± 0.007	0.713 ± 0.007	-0.014

Robustness to regularization. We sweep DropEdge rates (0%–40%) for both DELTA-Matched and GraphGPS (Appendix B). DELTA leads GraphGPS on 3p MRR at all five rates tested, with the advantage ranging from $+0.001$ (at 20%, within noise) to $+0.029$ (at 0%). The advantage is statistically meaningful at four of the five rates; the multi-hop advantage is structural, not an artifact of a particular hyperparameter configuration.

4.4 Standard Link Prediction

Table 6 shows standard link prediction results. GraphGPS leads on 1p MRR (0.513 vs. DELTA-Matched's 0.495), a gap of -0.018. However, DELTA-Matched achieves 97% of GraphGPS's MRR with only 69% of the parameters (158K vs. 228K). All models on the top-500 subgraph exceed published full FB15k-237 results for DistMult (0.241) and CompGCN (0.355) by large margins, reflecting the higher link density of the degree-based subgraph.

The LP vs. multi-hop contrast illustrates DELTA's architectural thesis: standard 1p evaluation rewards memorization of local link statistics, where GraphGPS's larger parameter budget provides an advantage. Multi-hop evaluation rewards compositional generalization, where DELTA's edge-to-edge attention provides a structural advantage that grows with depth.

4.5 Ablation: 1-Hop vs. 2-Hop Edge Adjacency

To directly validate the role of 2-hop edge adjacency—the architectural mechanism underlying DELTA's compositional reasoning—we ablate three configurations on the FB15k-237 subgraph using the multi-hop evaluation protocol:

- **1-hop**: edges attend only to directly adjacent edges (share a node)
- **2-hop**: edges attend to edges two steps away (default DELTA)
- **Node-only**: edge attention stream disabled; DELTA reduces to GAT-style node attention

Table 7 shows that on the top-500 subgraph all three configurations are statistically indistinguishable: every pairwise gap is within one standard deviation (3 seeds). Notably, the *node-only* condition (edge attention disabled) achieves the highest point estimates on every metric, though again within noise. This contrasts sharply with the synthetic transitive benchmark (2-hop 100% vs. 1-hop 61.1%), where compositional structure is unambiguous.

We interpret this as evidence that on this degree-biased subgraph, multi-hop MRR is dominated by entity embedding quality learned during link-prediction training; the edge adjacency mechanism provides a secondary contribution. The subgraph's high density (mean node degree ≈ 19.7) means 1-hop adjacency already covers nearly all structurally relevant edges. This motivates evaluating on the sparse full graph, where 1-hop neighborhoods are genuinely limited.

Table 6: Standard link prediction (1p) on FB15k-237 top-500 subgraph. Best-validation checkpoint, seed=1.

Model	**Params**	**MRR**	**H@1**	**H@3**	**H@10**	**Peak Ep**
GraphGPS	228K	**0.513**	**0.375**	0.581	0.813	200
SBHybrid	381K	0.509	0.363	**0.587**	**0.816**	250
DELTA-Matched	158K	0.495	0.346	0.572	0.804	200
DELTA-Full	293K	0.494	0.355	0.559	0.792	200
DistMult	47K	0.484	0.346	0.548	0.763	500+
GRIT	197K	0.439	0.295	0.496	0.760	200

Table 7: Hop-depth ablation (Phase 66): DELTA-Matched on FB15k-237 top-500 subgraph, 3 seeds $\times$ 500 epochs. Values are mean $\pm$ std. All differences are within 1σ; no condition is statistically superior.

Config	**LP MRR**	**1p MRR**	**2p MRR**	**3p MRR**	**2p$\rightarrow$3p**
2-hop (DELTA)	0.496 ± 0.003	0.534 ± 0.008	0.726 ± 0.014	0.731 ± 0.006	$+0.005$
1-hop	0.498 ± 0.008	0.551 ± 0.001	0.721 ± 0.007	0.729 ± 0.007	$+0.008$
Node-only	0.505 ± 0.006	0.553 ± 0.006	0.728 ± 0.011	0.742 ± 0.012	$+0.014$

Full-Scale Validation on Complete FB15k-237. To determine whether 2-hop edge adjacency provides measurable benefit on a sparse graph (mean degree ≈ 4.1), we repeat the 1-hop vs. 2-hop ablation on the full 14,541-entity FB15k-237 graph, using top-k=128 sparse edge attention to manage the 29.8M (hops=1) and 34.8M (hops=2) adjacency pairs. We train DELTA-Matched at full-graph scale (1,064K parameters, increased from 158K to accommodate 14,541 entities) for 200 epochs across 3 seeds per condition (Table 8).

Key findings from the full-scale ablation:

1. **2-hop improves multi-hop chain queries on the sparse graph.** Hops=2 provides $+0.003$ on 2p MRR and $+0.012$ on 3p MRR relative to hops=1 (final 3-seed averages). (Note: preliminary 2-seed results showed a larger $+0.017$ gap; seed 456's hops=2 3p score fell below the other seeds, revising the 3-seed average to $+0.012$, which still exceeds the pre-specified threshold.) The 3p improvement exceeds our pre-specified 0.010 meaningfulness threshold, directly confirming that 2-hop adjacency adds discriminative value when 1-hop neighborhoods are genuinely sparse. The 2p gap ($+0.003$) falls below the threshold, consistent with 2-hop adjacency providing incremental benefit primarily for longer reasoning chains.
2. **Overall LP MRR is statistically equivalent.** The two conditions achieve 0.205 and 0.204 LP MRR respectively—well within one standard deviation. As on the subgraph, 2-hop adjacency adds multi-hop benefit without harming link prediction.
3. **Full-graph LP performance is below published SOTA.** At 1p MRR ≈ 0.257, DELTA on full FB15k-237 is substantially below RotatE (0.338), CompGCN (0.355), and NBFNet (0.415). This reflects our 1-layer architecture, parameter scale constraints (DELTA was designed for compositional evaluation, not optimal 1p performance), and the absence of specialized KGC training strategies (e.g., self-adversarial negative sampling, N-N scoring). The multi-hop improvements reported are relative to our own hops=1 DELTA baseline under identical training conditions; they do not constitute a SOTA claim on full FB15k-237. Closing the gap to SOTA single-hop performance while maintaining the compositional advantage is an important open direction (§5).

Per-seed results (hops=1): All three seeds (42, 123, 456) fall within MRR 0.203–0.207, 2p MRR 0.109–0.117, and 3p MRR 0.203–0.214, confirming low cross-seed variance under the full-graph training regime. **Per-seed results (hops=2):** All three seeds complete: seed 42 achieves 2p MRR 0.119, 3p MRR 0.227; seed 123 achieves 2p MRR 0.118, 3p MRR 0.219; seed 456 achieves 2p MRR 0.111, 3p MRR 0.208. Seeds 42 and 123 are consistently above their hops=1 counterparts; seed 456 is comparable. Table 8 reports aggregate averages across all three seeds.

Table 8: Phase 67 — Full FB15k-237 hop-depth ablation (3 seeds $\times$ 200 epochs, top-k=128 sparse attention, 1,064K parameters). Values are mean $\pm$ std across all 3 seeds. 2p/3p queries use 10,000 samples each.

Config	**LP MRR**	**H@10**	**1p MRR**	**2p MRR**	**3p MRR**	**2p→3p**[†]
hops=1	0.205 ± 0.002	0.381 ± 0.003	0.257 ± 0.004	0.113 ± 0.003	0.206 ± 0.005	$+0.093$
hops=2	0.204 ± 0.002	0.382 ± 0.002	0.252 ± 0.003	0.116 ± 0.003	0.218 ± 0.008	$+0.102$
Δ (2−1)	-0.001	$+0.001$	-0.005	$+0.003$	$\mathbf{+0.012}$	—

[†]Absolute 2p→3p improvement (3p MRR − 2p MRR). Both conditions show positive improvement from 2p to 3p on the full graph, consistent with the subgraph finding.

Table 9: Phase 68 — Random-subgraph hop-depth ablation (3 seeds $\times$ 500 epochs, N=2,242 uniformly sampled entities, mean degree 7.6, hops=2 adjacency capped at 964K pairs). Values are mean $\pm$ std.

Config	**LP MRR**	**1p MRR**	**2p MRR**	**3p MRR**	**2p→3p**
hops=1	0.229 ± 0.027	0.334 ± 0.040	0.242 ± 0.057	0.274 ± 0.060	$+0.032$
hops=2	0.233 ± 0.019	0.340 ± 0.017	0.249 ± 0.028	0.284 ± 0.031	$+0.035$
Δ (2−1)	$+0.004$	$+0.006$	$+0.007$	$+0.010$	—

Density Control: Random-Subgraph Replication (Phase 68). Phase 66 used top-degree entities (mean degree 19.7), which the reviewer correctly identified as a potential confound: dense neighborhoods may inflate the observed pattern regardless of hop depth. To isolate the model property, we re-ran the hops=1 vs. hops=2 ablation on a *random* 2,242-entity subgraph of FB15k-237 (uniformly sampled, subgraph seed fixed to 0). We targeted N=2,500 uniformly sampled entities; 2,242 survived connectivity filtering, yielding mean degree 7.6—intermediate between Phase 67's sparse full graph ($\approx$4.1) and Phase 66's dense subgraph ($\approx$19.7), and trainable on a single RTX 3080 Ti in under three hours. Hops=2 adjacency pairs were subsampled to 964K to match the hops=1 count, providing an equal-attention-budget comparison (Table 9).

The between-conditions 3p gap is $+0.010$—larger than Phase 66's $+0.002$ but smaller than Phase 67's $+0.012$ (Table 8). Table 10 presents the cross-density picture at a glance.

Across all three density regimes, the depth-monotonic 2p→3p pattern persists under both conditions, confirming it is a property of query difficulty structure rather than graph density. The decreasing hops=2 advantage with increasing density is consistent with 1-hop neighborhoods becoming informationally sufficient above a density threshold.

4.6 Inference Efficiency

Table 11 reveals an important nuance in DELTA's computational profile. Encoding—a one-time graph processing step that includes 2-hop edge adjacency construction—is 51.8$\times$ slower than GraphGPS. However, *per-query scoring* is 0.84–0.94$\times$ GraphGPS: DELTA is faster per query. In deployment scenarios where the graph is encoded once and queries are scored many times, DELTA's amortized cost is comparable to or better than GraphGPS. Training cost (34.2$\times$) is a one-time expense dominated by edge adjacency computation.

4.7 Attention Temperature Asymmetry

We introduce learnable per-head temperature parameters $\tau^{(h)} = \exp(\theta^{(h)})$ that scale attention scores before softmax (Eq. 2). Higher τ produces sharper (more selective) attention; lower τ produces softer (more uniform) attention. Starting from uniform initialization, we observe how the model discovers differential sharpness preferences through gradient descent. We report these as observational findings; causal attribution would require controlled interventions such as temperature freezing, which we leave to future work.

Table 10: Cross-density comparison: hops=2 vs. hops=1 on 3p MRR across three density regimes (Phases 66–68). The 3p gap decreases monotonically with mean degree, consistent with 1-hop neighborhoods becoming informationally sufficient above a density threshold. Absolute 3p values are not directly comparable across rows (different entity pools and query difficulties).

Phase	**Subgraph**	**Mean Deg**	**hops=1 3p**	**hops=2 3p**	**Δ 3p**
67	Full FB15k-237 (N=14,541)	4.1	0.206 ± 0.005	0.218 ± 0.008	$\mathbf{+0.012}$
68	Random (N=2,242)	7.6	0.274 ± 0.060	0.284 ± 0.031	$+0.010$
66	Top-degree (N=500)	19.7	0.729 ± 0.007	0.731 ± 0.006	$+0.002$

Table 11: Inference timing comparison (mean of 10 runs, CUDA-synchronized).

Metric	**DELTA-Matched**	**GraphGPS**	**Ratio**
Encoding (once)	454.1 ms	8.8 ms	51.8×
1p per-query	777.7 μs	921.9 μs	**0.84**×
2p per-query	1,380 μs	1,476 μs	**0.94**×
3p per-query	1,251 μs	1,371 μs	**0.91**×
Training (total)	3,782 s	110 s	34.2×

Edge–node temperature divergence. Across all configurations tested (20+ over 9 experimental phases), a consistent pattern emerges: **edge attention temperatures drift upward** (from 4.0 to 4.4–5.0) while **node attention temperatures drift downward** (from 4.0 to 3.5–3.7). The model autonomously discovers that edges benefit from selective, discriminative attention (identifying which relational compositions matter), while nodes benefit from broader, averaging attention (aggregating diverse neighborhood signals).

Layer 0 is always dead. Regardless of temperature initialization, first-layer attention heads remain near-uniform (entropy > 0.95). This is architecturally expected: initial entity embeddings carry no relational information worth selectively attending to. Useful attention emerges only after the first reconciliation bridge provides cross-stream context. This observation has an important architectural implication: if first-layer attention is uniformly near-random, the reconciliation bridge—not the edge-to-edge attention mechanism itself—may initiate the representational differentiation that later layers exploit. Whether the bridge's sequential cascade (edges absorbing node context first, then nodes absorbing the updated edge context) is the primary load-bearing mechanism, rather than the edge-to-edge attention it conditions, is an open question that a direct bridge-ablation experiment could resolve.

Temperature controls operating mode. We identify three distinct operating regimes:

1. **LP-optimized** (asymmetric: node=2, edge=6): LP MRR 0.491, sacrifices multi-hop.
2. **Balanced** (uniform temp=4.0): Best 3p MRR (0.402), moderate LP.
3. **Annealed** (node temp $4.0 \rightarrow 2.0$ during training): Breaks the 3p ceiling (0.415), approaching combined LP+3p optimality.

Temperature annealing from high (exploration) to low (exploitation) during training creates compositional representations in early epochs that static initialization cannot replicate. This finding—that training *trajectory* matters, not just final temperature—connects to curriculum learning and suggests that compositional representations require a specific developmental sequence.

LP improvements are robust; multi-hop temperature effects are not. Multi-seed validation (3 seeds) confirms that temperature-driven LP improvements are statistically robust (MRR 0.484 ± 0.005). However, multi-hop MRR differences between temperature configurations showed high cross-seed variance. Increasing query counts to 10,000 reduced cross-seed std by 66–84%, confirming that evaluation protocol sensitivity was the dominant variance source. We do not make strong claims about temperature effects on multi-hop reasoning; the primary validated effect is on link prediction.

Table 12: Scaling analysis: DELTA (1-layer) vs. DistMult across entity counts. All models on FB15k-237 subgraphs, seed=42. The full FB15k-237 row (Phase 67) uses top-k=128 sparse attention and 1,064K parameters; DM result is from published literature.

N	**Triples**	$\lvert E_{\text{adj}}\rvert$	**Test MRR**		Δ	**DELTA time**	**DM time**	
			DELTA	DM				
500	9,700	256K	0.495	0.484	+0.011	3,782s	110s	†hops=1
2,000	62,700	15.2M	0.334	0.319	+0.015	5,788s	72s	
5,000	152,800	63.0M	0.247	0.224	+0.023	21,688s	1,784s	
14,541	310,116	29.8M†	0.205	0.241*	n/a*	—	—	

adjacency. *__Not directly comparable:__ the published DistMult result [2] uses tuned hyperparameters and extended training time; our DELTA result at N=14,541 uses a fixed training budget. The $\Delta=-0.036$ entry does *not* establish that DELTA underperforms DistMult under matched conditions at full scale; see §4.8 for discussion.

4.8 Scaling Analysis

We evaluate DELTA's scaling behavior across graph sizes from N=500 to the full N=14,541 entity FB15k-237 graph (Table 12).

Genuine but modest advantage at scale. DELTA consistently outperforms DistMult at all controlled scales tested (N=500 to N=5,000), with the gap growing from $+0.011$ to $+0.023$. At full FB15k-237 (N=14,541), DELTA achieves MRR 0.205 vs. published DistMult 0.241; this gap reflects that DistMult at full scale benefits from extensive hyperparameter tuning and training time in the original publication, while our DELTA result uses a fixed training budget. A fair comparison between our DELTA and a freshly trained DistMult baseline on the full graph would require matched training conditions and is left to future work.

Depth is the scaling bottleneck, not the mechanism. At N=2,000 with 15.2M edge adjacency pairs, 3-layer DELTA catastrophically over-smooths (MRR 0.002, near random), while 1-layer DELTA surpasses DistMult. Each additional layer compounds representation homogenization through the dense adjacency graph. Residual gating (Phase 60) eliminates the failure mode but gates freeze at $\sim$10% layer / 90% residual, confirming that additional depth provides no benefit at this scale.

Attention dilution is the primary scaling challenge. At N=5,000, each edge has hundreds of adjacency neighbors. Full softmax over these large neighborhoods produces near-uniform attention weights, requiring $\sim$50 training epochs to learn discriminative patterns (vs. $\sim$25 epochs at N=500). Subsampling the edge adjacency to $\sim$48% retention slightly *improves* MRR ($+0.007$), suggesting that a moderate number of high-quality neighbors outperforms access to all neighbors.

Top-k sparse attention is a validated scaling mechanism. At k=128, sparse attention on the full 63M adjacency pairs matches full-softmax quality exactly (test MRR 0.2472 vs. 0.2457), while k=64 degrades by -5.6%. The quality cliff lies between 64 and 128 neighbors. This validates a path to larger-scale deployment: for graphs where full-softmax edge attention is memory-prohibitive, k=128 provides lossless approximation.

5 Discussion

Why does the advantage grow with depth? We hypothesize that 2-hop edge adjacency creates a multiplicative advantage for compositional reasoning. At each hop, DELTA composes relational information directly in edge space, while node-based models must decompress through node representations. For single-hop predictions, this overhead is negligible—node-level representations suffice. But at 5 hops, the accumulated compression loss in node-mediated architectures becomes substantial, while DELTA's direct edge-to-edge channel preserves compositional fidelity. The temperature analysis supports this interpretation: edge attention's preference for sharpness indicates selective relational composition, not broad averaging.

Subgraph vs. full-graph: when does the mechanism matter? The Phase 66 subgraph ablation found all adjacency conditions statistically indistinguishable (mean degree ≈ 19.7), while the Phase 67 full-graph ablation found a meaningful 3p MRR gain of $+0.012$ (mean degree ≈ 4.1). This is consistent with the mechanism's design rationale: on dense graphs, 1-hop adjacency already reaches the structurally relevant edges a relational composition would need; on sparse graphs, 2-hop adjacency provides a genuinely expanded receptive field. The mechanism is most valuable when graph sparsity creates compositional gaps that 1-hop cannot bridge.

Why does capacity hurt multi-hop? DELTA-Matched (158K) consistently outperforms DELTA-Full (293K) on multi-hop tasks despite having fewer parameters. We attribute this to an overfitting–generalization trade-off: larger models have sufficient capacity to memorize local link statistics (achieving higher 1p MRR), which provides a shortcut that bypasses compositional reasoning. The constrained model must learn genuinely compositional representations because it cannot afford to memorize.

Does the compositional inductive bias trade off against 1p performance? A legitimate concern is that the architectural choices enabling DELTA's structural inductive bias are themselves in tension with competitive single-hop LP. Two observations bear on this. First, the LP/multi-hop gap is stable across scales: at $N{=}500$, DELTA trails GraphGPS by -0.018 on 1p while leading by $+0.100$ at 5p; the compositional advantage does not cost proportionally more LP as scale increases. Second, DELTA's full-graph LP gap below SOTA (0.205 vs. 0.338–0.415) is far larger than the -0.018 gap at $N{=}500$, suggesting the full-graph deficit is primarily a training scale and strategy issue (self-adversarial negative sampling, multi-layer training with gating, longer schedules) rather than an architectural incompatibility. We cannot directly verify this: whether a SOTA-performing DELTA variant preserves the depth-monotonic compositional advantage remains the primary open question for future work.

Limitations. We enumerate limitations explicitly, organized by severity.

Scale and benchmark coverage (critical). Our compositional reasoning experiments use a 500-entity degree-biased subgraph of FB15k-237, which has substantially higher link density (mean degree ≈ 19.7) than the full graph (mean degree ≈ 4.1). The strong multi-hop results (0.758–0.790 MRR, 5-hop) are demonstrated on this small subgraph. On the complete 14,541-entity FB15k-237 graph, DELTA achieves 1p MRR ≈ 0.205, substantially below published SOTA: RotatE 0.338, CompGCN 0.355, NBFNet 0.415. We have not evaluated on other standard benchmarks (YAGO3-10, WN18RR [24], ogbl-wikikg2), which is an important open direction.

Baseline coverage (important). We compare primarily against general graph transformers (GraphGPS, GRIT) and the DistMult embedding baseline. We do not include specialized path-based KG reasoners (NBFNet [8], A*Net [21]), which explicitly model relational paths and are more directly comparable to DELTA's multi-hop reasoning claim. Including these is a high-priority next step. Our multi-hop evaluation uses internally generated chain queries with soft entity traversal—not the established BetaE/ConE/GQE complex-query benchmarks, which use different query generation protocols, logical operators (conjunction, disjunction), and evaluation metrics. Results are not directly comparable to published multi-hop reasoning numbers from those systems.

Full-scale multi-hop evaluation (important). The Phase 67 full-graph ablation evaluates hops=1 vs. hops=2 on standard link prediction MRR and 2p/3p chain query MRR, but does not include the full 2-hop$\rightarrow$5-hop depth sweep reported in the subgraph experiments. Whether DELTA's depth-monotonic improvement persists on the full sparse graph has not been directly validated.

Computational cost (notable). The one-time encoding overhead (51.8$\times$ GraphGPS) limits applicability to static graphs. Training cost (34.2$\times$ GraphGPS at $N{=}500$) may be prohibitive for practitioners with limited compute budgets. The full-graph (Phase 67) experiment required approximately 3 days on an NVIDIA H200 GPU with 141 GB VRAM to evaluate 6 conditions (3 seeds $\times$ 2 hop settings $\times$ 200 epochs), with seeds running in parallel within each condition.

Causal claims (methodological). The attention temperature asymmetry is reported as a consistent observational pattern across 20+ configurations, not as a causally established property. Causal attribution would require controlled interventions (temperature freezing, ablations of individual head types) that we have not performed. The brain architecture (Appendix C) and domain transfer

(Appendix D) results are preliminary signals from small experiments and should not be treated as validated findings.

Subgraph selection bias (methodological). Using a top-degree subgraph for primary experiments biases toward entities with rich relational context, which may inflate multi-hop performance relative to a random subgraph or the full graph. Our subgraph's density also means the edge adjacency graph is much more saturated than on sparse realistic graphs, which may affect which architectural choices appear beneficial.

6 Conclusion

We have presented DELTA, an architecture that treats edges as first-class computational entities via edge-to-edge attention over multi-hop adjacency. Across three complementary experimental contexts—synthetic structural benchmarks (5 seeds, cross-task), a full-scale 14,541-entity graph ablation (3 seeds), and a degree-based subgraph characterization (7 models, 5 depths)—a consistent picture emerges: edge-to-edge attention provides a structural reasoning advantage that grows with compositional depth and is most discriminative when graph sparsity creates chains that 1-hop adjacency cannot bridge. On the synthetic benchmarks, DELTA outperforms GraphGPS and GRIT by $+0.573$, $+0.290$, and $+0.095$ accuracy—the clearest cross-task evidence for the mechanism. On the complete FB15k-237 graph, 2-hop adjacency improves 3-hop chain MRR by $+0.012$ over 1-hop, confirming the mechanism at realistic scale. On the degree-based subgraph, DELTA-Matched is the only model with a depth-monotonic MRR improvement from 2-hop to 5-hop (GraphGPS: -0.064), though a direct ablation on the same dense subgraph finds hop-depth conditions statistically indistinguishable—consistent with the mechanism mattering less where 1-hop adjacency already saturates the structural neighborhood. Mechanistic analysis reveals an emergent edge–node temperature asymmetry providing interpretable evidence for the distinct computational roles of the two attention streams.

These results suggest that the standard single-hop link prediction benchmark undervalues architectures with strong relational inductive biases. Multi-hop evaluation provides a more discriminating test of compositional capacity, and we encourage its adoption alongside standard metrics. We note clearly that DELTA's full-graph LP MRR (≈ 0.205) falls below published SOTA methods (≈ 0.34–0.41), reflecting the gap between demonstrating a novel mechanism and achieving optimized end-to-end performance—a gap that remains an important open problem.

Future work includes (1) full-depth compositional evaluation (2-hop→5-hop) on the complete 14,541-entity graph; (2) comparison against path-based baselines NBFNet [8] and A*Net [21]; (3) evaluation on established complex query benchmarks (BetaE/GQE query splits [15]) to enable direct comparison with specialized multi-hop reasoners; (4) cross-dataset validation on YAGO3-10, WN18RR, and ogbl-wikikg2; (5) investigation of the BrainEncoder differentiable graph construction module (directional preliminary signal only: $+4.7\%$ Hits@10 from a 100-sample probe with $\pm 5\%$ confidence intervals—not a validated finding, reported for completeness in Appendix C); and (6) closing the gap to SOTA single-hop LP performance by incorporating specialized KGC training strategies while preserving the multi-hop compositional advantage. The top-k sparse attention mechanism validated at N=5,000–14,541 provides a concrete path to scaling edge-to-edge attention to large knowledge graphs.

References

[1] Antoine Bordes, Nicolas Usunier, Alberto Garcia-Duran, Jason Weston, and Oksana Yakhnenko. Translating embeddings for modeling multi-relational data. In *NeurIPS*, 2013.

[2] Bishan Yang, Wen-tau Yih, Xiaodong He, Jianfeng Gao, and Li Deng. Embedding entities and relations for learning and inference in knowledge bases. In *ICLR*, 2015.

[3] Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. RotatE: Knowledge graph embedding by relational rotation in complex space. In *ICLR*, 2019.

[4] Théo Trouillon, Johannes Welbl, Sebastian Riedel, Éric Gaussier, and Guillaume Bouchard. Complex embeddings for simple link prediction. In *ICML*, 2016.

[5] Linlin Chao, Jianshan He, Taifeng Wang, and Wei Chu. PairRE: Knowledge graph embeddings via paired relation vectors. In *ACL*, 2021.

[6] Michael Schlichtkrull, Thomas N. Kipf, Peter Bloem, Rianne van den Berg, Ivan Titov, and Max Welling. Modeling relational data with graph convolutional networks. In *ESWC*, 2018.

[7] Shikhar Vashishth, Soumya Sanyal, Vikram Nitin, and Partha Talukdar. Composition-based multi-relational graph convolutional networks. In *ICLR*, 2020.

[8] Zhaocheng Zhu, Zuobai Zhang, Louis-Pascal Xhonneux, and Jian Tang. Neural Bellman-Ford networks: A general graph neural network framework for link prediction. In *NeurIPS*, 2021.

[9] Yongqi Zhang and Quanming Yao. Knowledge graph reasoning with relational digraph. In *NeurIPS*, 2022.

[10] Ladislav Rampášek, Michael Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, and Dominique Beaini. Recipe for a general, powerful, scalable graph transformer. In *NeurIPS*, 2022.

[11] Liheng Ma, Chen Lin, Derek Lim, Adriana Romero-Soriano, Puneet K. Dokania, Mark Coates, Philip Torr, and Ser-Nam Lim. Graph inductive biases in transformers without message passing. In *ICML*, 2023.

[12] Hamed Shirzad, Ameya Velingker, Balaji Venkatachalam, Danica J. Sutherland, and Ali Kemal Sinop. Exphormer: Sparse transformers for graphs. In *ICML*, 2023.

[13] Zhengdao Chen, Xiang Li, and Joan Bruna. Supervised community detection with line graph neural networks. In *ICLR*, 2019.

[14] Jiaqi Gong, Shichao Pei, Jiaxin Bai, and Xiangliang Zhang. EHGNN: Edge heterogeneous graph neural network. *arXiv preprint arXiv:2310.07242*, 2023.

[15] Hongyu Ren and Jure Leskovec. Beta embeddings for multi-hop logical reasoning in knowledge graphs. In *NeurIPS*, 2020.

[16] Zhanqiu Zhang, Jie Wang, Jiajun Chen, Shuiwang Ji, and Feng Wu. ConE: Cone embeddings for multi-hop reasoning over knowledge graphs. In *NeurIPS*, 2021.

[17] William L. Hamilton, Payal Bajaj, Marinka Zitnik, Dan Jurafsky, and Jure Leskovec. Embedding logical queries on knowledge graphs. In *NeurIPS*, 2018.

[18] Hongyu Ren, Weihua Hu, and Jure Leskovec. Query2Box: Reasoning over knowledge graphs in vector space using box embeddings. In *ICLR*, 2020.

[19] Kristina Toutanova and Danqi Chen. Observed versus latent features for knowledge base and text inference. In *3rd Workshop on Continuous Vector Space Models and their Compositionality*, 2015.

[20] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *ICLR*, 2018.

[21] Zhaocheng Zhu, Xinyu Yuan, Michael Galkin, Sophie Xhonneux, Ming Zhang, Maxime Gazeau, and Jian Tang. A*Net: A scalable path-based reasoning approach for knowledge graphs. In *NeurIPS*, 2023.

[22] Erik Arakelyan, Daniel Daza, Pasquale Minervini, and Michael Cochez. Complex query answering with neural link predictors. In *ICLR*, 2021.

[23] Zhaocheng Zhu, Mikhail Galkin, Zuobai Zhang, and Jian Tang. Neural-symbolic models for logical queries on knowledge graphs. In *ICML*, 2022.

[24] Tim Dettmers, Pasquale Minervini, Pontus Stenetorp, and Sebastian Riedel. Convolutional 2D knowledge graph embeddings. In *AAAI*, 2018.

A Hits@10 by Depth

Table 13: Hits@10 across reasoning depths (single seed, same experiment as Table 3).

Model	1p	2p	3p	4p	5p
DELTA-Matched	0.854	0.889	**0.881**	**0.885**	**0.895**
GraphGPS	**0.848**	**0.897**	0.859	0.818	0.826
DistMult	0.800	0.858	0.744	0.672	0.587

B DropEdge Robustness

Table 14: 3-hop MRR across DropEdge rates. DELTA-Matched leads at all rates.

Model	0%	10%	20%	30%	40%
DELTA-Matched	0.740	0.744	0.724	0.732	0.744
GraphGPS	0.711	0.716	0.723	0.720	0.725
Δ	+0.029	+0.028	+0.001	+0.012	+0.019

C Brain Architecture (Preliminary)

The BrainEncoder module uses Gumbel-sigmoid differentiable edge selection to construct new edges from learned node representations, then processes the augmented graph (original + constructed edges) through DELTA layers. On the FB15k-237 subgraph (N=494):

Table 15: BrainEncoder results (multi-seed, density d=0.01, 200 epochs).

Config	MRR	H@1	H@10	Constructed edges
DELTA-Full	0.494	0.355	0.760	0
brain_hybrid (d=0.01)	0.484 ± 0.010	0.334 ± 0.011	$\mathbf{0.799 \pm 0.006}$	2,435

The brain_hybrid configuration matches DELTA-Full on MRR while delivering $+4.7\%$ Hits@10, demonstrating that differentiable graph construction adds genuine recall through learned structural augmentation. Density is a critical hyperparameter: d=0.01 (2,435 edges) strictly dominates d=0.02 (4,870 edges) on all metrics.

D Domain Transfer (Preliminary)

A frozen DELTA encoder trained on FB15k-237, evaluated via a linear probe on WN18RR, achieves accuracy of **0.961**—without any domain adaptation mechanism (no GRL, no fine-tuning). **Note:** this evaluation uses 100 training samples, giving approximate $\pm$5% confidence intervals; results should be treated as a preliminary signal, not a statistically validated finding. A proper transfer evaluation requires $\geq$10,000 samples with confidence intervals on held-out test splits. We report this result as a directional indicator that DELTA's learned edge representations may capture transferable relational structure.